

Compiler Project

Code Generation for VSOP

Roland Greffe & Pascal Fontaine

University of Liège

April 14, 2026

Outline

The Assignment

The LLVM Language

Code Generation

Implementing the Dispatch

Q & A

Outline

The Assignment


The LLVM Language

Code Generation

Implementing the Dispatch

Q & A

Code Generation



For **semantically valid** VSOP programs, generate **LLVM IR**

No error should be found during this phase

`vsopc -i example.vso`p should:

- If the source file is valid, return 0 and dump on stdout the LLVM IR
- Otherwise, return a non-zero value and print on stderr (at least) one error

Generating an Executable

`vsopc /path/example.vsop` should:

- If the source file is valid, return 0 and create an executable `/path/example` from the generated IR
 - This executable will then be run without argument
- Otherwise, return a non-zero value and print on `stderr` (at least) one error

Assignment

Due the 13th of May

Automated tests worth 5% of your grade

You can use **LLVM bindings** to generate the IR, an example is provided on eCampus

Support for **custom tests** in tests subfolder: to **give inputs** to the generated executable, provide a `/path/example.input` file (same name that the source file but with the `input` extension) which will be used as `stdin`

Five modes: no flag, `-i`, `-c`, `-p` and `-l`

Assignment

Optional: you can implement some **VSOP extensions** of your liking

- But they should **only be enabled** when `vsopc` is called with the `-e`
- The `-e` flag should be able to be used with the other flags

A (short) **report** should be provided with your code discussing:

- the **implementation** of your compiler
- the potential **shift/reduce** or **reduce/reduce conflicts** you had to solve
- the **potential extensions** you have implemented
- ...

Output Format

When called with the **-i flag**, your compiler should dump on stdout the **generated IR**

When called with **no flag**, **nothing** should be printed on stdout

Error Management

Your compiler should not output any error at code generation:

All errors (that could have been detected at compile-time) **should already have been reported**

Of course, it will still report potential **lexical**, **syntax** or **semantic** errors

Report

Provide a **PDF report** describing:

- How your code is organised, which tools, data structures and algorithms are used
- The potential **shift/reduce or reduce/reduce conflicts** you had to solve
- If something in your implementation is not obvious from your **documented** code
- Your **VSOP extensions**
- The **limitations** of your compiler: what would you do differently or with additional time ?
- How long did it take, what were the main difficulties ? How can we improve the project ?

Try to be **succinct**

Outline

The Assignment

The LLVM Language

Code Generation

Implementing the Dispatch

Q & A

What is LLVM ?

- Originally, acronym for *Low Level Virtual Machine*
 - But since its creation, the LLVM project has well evolved
 - Nowadays, **little relationship** with *virtual machines*
- Collection of **modular and reusable** tools to support both **static and dynamic** compilation of arbitrary programming languages
- Many sub-projects, the most prominent one being **clang**, a C/C++ compiler
- Wildly used in the industry (e.g. Adobe, Apple, Cray, Intel, NVIDIA, Siemens, Sun, ...).
- Somewhat **easy, fast, modular**, and **supports many targets**.

What is the LLVM Language

- An *Intermediate Representation* (**IR**)
- **Assembly**-like language
- Based on the *Single Static Assignment* (**SSA**) form
- **Typed** !
- Allows **low-level** operations, while permitting to represent **high-level** languages cleanly (e.g., function definitions)

Resources

- LLVM Language **Reference Manual** ([Link](#))
- The Often Misunderstood **GEP** Instruction ([Link](#))
- The LLVM **Tutorial**: Kaleidoscope: Implementing a Language with LLVM ([Link](#))
- Static Single Assignment (SSA) Form (Wikipedia) ([Link](#))
- The (**not so**) theoretical course !
- Mapping **High-Level Constructs** to LLVM IR ([Link](#))

Practical Details

Version: **Use LLVM version 11** available in the reference container

How to **generate LLVM code**:

- **Textually**, then call LLVM assembler (and optimizer)
- Using the **LLVM bindings library**:
 - Easier to optimize
 - Allows interpretation and JIT compilation
 - **Poorly documented**, steep learning curve !
 - Official bindings for C/C++ and various unofficial bindings for other languages (e.g., `llvmlite` for Python)
 - Example for C/C++ on eCampus

LLVM Modules

A compilation unit is called a **module**, it contains:

- **Type aliases**
- **Declarations**, which states that a symbol exists and give its type
- Definitions of **global variables** (type + initializer)
- **Function definitions**
- Special sections such as module initialization code

A module lies in its own file (`module_name.ll`) or is generated in-memory using the library

You are expected to generate a **single module** in this project, but adding the possibility to import external modules is a possible extension

LLVM Types

<code>i1</code>	<i>; 1-bit integer (may be used for bool)</i>
<code>i8</code>	<i>; 8-bit integer (used for char)</i>
<code>i32</code>	<i>; 32-bit integer</i>
<code>i64</code>	<i>; 64-bit integer</i>
<code>float</code>	<i>; 32-bit IEEE 754 floating-point number</i>
<code>double</code>	<i>; 64-bit IEEE 754 floating-point number</i>
<code>void</code>	<i>; empty type</i>
<code>label</code>	<i>; named addresses (see later)</i>
<code>i32 *</code>	<i>; pointer to one (or more) i32</i>
<code>[40 x i8]</code>	<i>; array of 40 8-bit numbers (static size)</i>
<code>float (i16, i32 *)</code>	<i>; function (i16, i32 *) -> float</i>
<code>i32 (i8 *, ...)</code>	<i>; signature of printf()</i>
<code>{ i32, float }</code>	<i>; structure with i32 and float fields</i>
<code>%MyVeryOwnType</code>	<i>; named type</i>

SSA Values

Also known as *registers* (somewhat improperly)

They are **immutable**, i.e. do not change after initialization

Two kinds of identifiers:

- **Global identifiers** (functions, global variables) begin with the @ character (e.g. @puts, @my_global_var)
- **Local identifiers** (register names) and types begin with the % character (e.g. %result, %MyStruct)

Infinite supply of registers and globals

Unnamed values are (and **must** be) **numbered sequentially** (e.g. %31, %32, %33, ...)

Literals

Simple literals:

- **Integers:** e.g. 42, -1028
- **Booleans:** `true` and `false` of `i1` type (resp. 1 and 0)
- **Floating-point:** usual notation (e.g. 3.14159265, 6.673981e-11), but requires **exact decimal value** (e.g. 0.1 is rejected), or direct representation (e.g. 0x141d7038)
- **Null pointer:** `null`, of any pointer type

Complex literals:

- **Structures:** e.g. { `i32` 4, `float` 17.0, `i32*` @g }
- **Arrays:** [`i32` 1, `i32` 2, `i32` 3]
- **Zeroes:** `zeroinitializer`
- **Undefined:** `undef`
- Global variables and functions are **always implicit pointers**:
 - i.e. @x always has a pointer type

Blocks, Labels and Terminator Instructions

A function body is a set of **SSA blocks**

A block spans from a **label** to a **terminating instruction**:

```
ret void                ; return from procedure
ret i32 42              ; return 42
ret { i32, double } { i32 42, double 84.0 } ; return struct literal
br label %end_of_process ; unconditional jump
br i1 %cond, label %if_true, label %if_false ; conditional jump
```

Example:

```
br i1 %cond, label %if_true, label %if_false
; A statement here would be illegal

if_true:                                ; Labels are as in C
ret i32 42

if_false:
ret i32 0
```

Arithmetic Instructions

Remarks:

- **Separate instructions** for **integers** and **floating-point** numbers (e.g. `add` and `fadd`)
- Operands type (and size) must be **known**
- No signedness in types, but **signedness in some instructions** (e.g. `udiv` and `sdiv`)

```
%1 = add i32 42, %var      ; integers addition, %var must also be i32
%3 = fadd double 42.0, %x  ; different instruction for floats
%4 = sub i32 0, %var       ; yields -%var
%5 = fmul double 3.0, %x   ; multiplication for floats
%6 = udiv i32 %var, 5       ; unsigned integers division
%7 = sdiv i32 %var, 5       ; signed integers division (rounded towards 0)
%8 = fdiv double %x, 5.0   ; floats division
```

There exist other operations (and flags), see full reference

Logical Instructions

Remark: the logical instructions are **always bitwise**

`%1 = and i1 %a, %b ; 'and' on two booleans`

`%2 = and i32 9, 10 ; bitwise, yields 8 (1001 & 1010 = 1000)`

`%3 = or i32 9, 10 ; yields 11 (1001 | 1010 = 1011)`

`%4 = xor i32 9, 10 ; yields 3 (1001 ^ 1010 = 0011)`

`%5 = xor i32 %x, -1 ; yields ~%x`


Comparisons

Remarks:

- Comparisons are made with the **icmp** (for integers) and the **fcmp** (for floats) instructions followed by a flag indicating the comparison type
- Comparisons always return an **i1**

```
%1 = icmp eq i32 %a, %b      ; %a == %b, returns i1 type
%2 = icmp ne i32 %a, %b      ; %a != %b
%3 = icmp ugt i32 %a, %b      ; unsigned %a > %b
%4 = icmp sle i32 %a, %b      ; signed %a <= %b
%5 = icmp eq i32* @p, @q      ; works also on pointers
%7 = fcmp oeq float %x, %y    ; ordered equal, i.e.
                                ; %x != NaN && %y != NaN && %x == %y
%8 = fcmp ueq float %x, %y    ; unordered equal, i.e.
                                ; %x == NaN || %y == NaN || %x == %y
%9 = fcmp olt float %x, %y    ; %x != NaN && %y != NaN && %x < %y
```

The Magical Phi Instruction



The Magical Phi Instruction in LLVM

phi takes on the value specified by the pair corresponding to the block that executed **just before** the current block

```
br i1 %cond, label %if_true, label %if_false
if_true:
    %i.1 = i32 42
    br label %end_if
if_false:
    %i.2 = i32 0
    br label %end_if
end_if:
    %i = phi i32 [%i.1, %if_true], [%i.2, %if_false]
```

Constants values are allowed as well

Functions Declarations and Calls

Declarations:

```
; use '...' to have variable number of arguments  
declare i32 @printf(i8*, ...)  
; use C (default) calling convention  
declare ccc i32 @printf(i8*, ...)  
; use fast calling convention  
declare fastcc void @someFastCCFun(i32, i8*)
```

Calls:

```
%z = call double @pow(double %x, double %y)  
call i32(i8*, ...)* @printf(i8* @fmt, i32 %val)  
call fastcc void @someFastCCFun(i32 42, i8* @str)
```

- Functions are **always pointers**
- Catching return value is not mandatory, but catching **void** is **illegal**

Memory Accesses

```
%ptr = alloca i32          ; yields i32* to STACK-allocated i32
store i32 42, i32* %ptr    ; stores 42 into allocated cell
%val = load i32, i32* %ptr ; %val := 42 (read from cell)
```

```
%aPtr = alloca i32, i64 100 ; int32_t aPtr[100];
; fill(aPtr, 100)
call fastcc void @fill(i32* %aPtr, i32 100)
; int *a4ptr = &aPtr[4];
%a4ptr = getelementptr i32, i32* %aPtr, i64 4
%a4 = load i32, i32* %a4ptr ; int a4 = *a4ptr;
```

```
%sPtr = alloca {i32, i8}    ; struct { int32_t a, int8_t b } s;
%bPtr = getelementptr {i32, i8}, {i32, i8}* %sPtr, i32 0, i32 1
store i8 42, i8* %bPtr      ; s.b = 42;
```

- Structures are indexed with **i32**, any integer for arrays
- Use `malloc` for heap-allocated memory

The `getelementptr` Instruction

This instruction is used to access an element from an array or a field of a structure, but **no memory dereferencing** is performed with this instruction

In C, **address computation** and **dereferencing** are performed within the **same instruction**:

```
int x = array[2]
int f = my_struct.f
```

While in LLVM, they are **separated**:

```
%x_addr = getelementptr i32, i32* %array, i64 2
%x = load i32, i32* %x_addr
%f_addr = getelementptr {i8, i32}, {i8, i32}* %my_struct, i32 0, i32 1
%f = load i32, i32* %f_addr
```

Note: 2 indexes are used when dealing with a structure, **the first one tells which structure to access** (only one possibility is this case, but what if we have an array of structures ?) and the second one **which field to access**

Function Definitions

```
define fastcc void @fill(i32* %array, i32 %len) {  
entry:  ; If not specified, label %0 will be inserted  
    br label %loop_cond
```

```
loop_cond:  
    %i = phi i32 [ 0, %entry], [ %ip1, %loop_body ]  
    %cond = icmp ult i32 %i, %len  
    br i1 %cond, label %loop_body, label %loop_end
```

```
loop_body:  
    %ptr = getelementptr i32, i32* %array, i32 %i  
    store i32 %i, i32* %ptr  
    %ip1 = add i32 %i, 1  
    br label %loop_cond
```

```
loop_end:  
    ret void
```

```
}
```

Translation of Lecture IR

Quite similar, but **typed** and **no assignments**

```
; LABEL my_label
```

```
my_label:
```

```
; GOTO my_label
```

```
br label %my_label
```

```
; id := 42      (7th assignment to id in that scope)
```

```
%id.6 = add i32 0, 42
```

```
; id := a + 1  (assuming i32 type)
```

```
%id = add i32 %a, 1
```

```
; id := g      (where g assigned in a if-then-else above)
```

```
%id = phi i32 [%g.1, %if_true], [%g.2, %if_false]
```

Translation of Lecture IR (cnt'd)

```
; id := M[addr]
```

```
%id = load i32, i32* %ptr
```

```
; M[addr] := id
```

```
store i32 %id, i32* %ptr
```

```
; a_exp_addr := exp * 4
```

```
; a_exp_addr := a_exp_addr + a_base_addr
```

```
%a_exp_ptr = getelementptr i32, i32* %a, i32 %exp
```

```
; IF n > 0 THEN if_greater ELSE if_lower
```

```
%cond = icmp sgt i32 %n, 0
```

```
br i1 %cond, label %if_greater, label %if_lower
```

```
; ret = CALL my_fun(arg1, arg2)
```

```
%val = call i32 @my_fun(i32 %arg1, i32 %arg2)
```

```
ret i32 %val
```

Obtain the Size of a Structure %T

Take the address of the second element of an hypothetical array of %T starting at address 0:

```
%size_as_ptr = getelementptr %T, %T* null, i32 1  
%size_as_i32 = ptrtoint %T* %size_as_ptr to i32
```

As `getelementptr` returns a pointer, you must do a **conversion** to an integer afterwards

Avoid trying to compute the size of a structure yourself: its size may be **target-dependent** because of **alignment issues**

Strings

In VSOP, all strings are ultimately string literals:

- They are **not mutable**
- They support **no operation** other than printing

LLVM support string literals **natively**, as arrays of `i8`:

```
@s = constant [14 x i8] c"Hello, world!\00"
```

```
declare i32 @puts(i8*)
```

```
define i32 @main() {  
    %1 = getelementptr [14 x i8], [14 x i8]* @s, i64 0, i64 0  
    %2 = call i32 @puts(i8* %1)  
    ret i32 0  
}
```

You can store them as `i8*`

What's Next

We haven't covered all of LLVM language (not even close)

Have a look at the language reference and the **tutorial**

Use `clang -S -emit-llvm` to get the LLVM IR when compiling with `clang`

Notable omissions:

- tail call optimization
- function and argument attributes
- conversion and casting operations
- LLVM intrinsics (e.g. for pow)
- linkage types
- undef and poison values
- vectors and vectored operations
- parallelism-oriented features
- memory management features
- special control flow (e.g. for exceptions)
- low-level stuff (e.g. access to volatile memory)

The LLVM Library

The LLVM Library allows:

- declarations of (external) functions and global variables
- definitions of global variables and functions
- adding blocks to functions
- moving the insertion point (*builder position*) at start/end of any block
- emitting instructions
- checking functions/modules are well-formed
- executing optimization passes in any order

However, it is:

- **poorly documented** (see the tutorial and Doxygen)
- not mandatory

A small **example** is available on eCampus

Compilation/Execution with the Library

Once you built (and optimized) a module, you can:

- **interpret** it using a built-in interpreter
- **execute** it after *just-in-time* (JIT) compilation
- **compile** it into a native executable
- **dumps** its LLVM assembly code

Outline

The Assignment

The LLVM Language

Code Generation

Implementing the Dispatch

Q & A


AST Traversal Reminder

You should **refer** to the (not so) theoretical course

Generally a mix between **top-down** and **bottom-up** approaches

At each node:

- 1 Receive some info from parent node (inherited)
- 2 Update received info based on current node
- 3 Process first child, passing needed info
- 4 Get back info from child processing (synth.)
- 5 Process that info
- 6 Process next child, passing needed info
- 7 Get back info from child processing
- 8 Process that info
- 9 Return useful info back to parent node



Do as many passes as needed

Factorize Common Code

VSOP being an **expression-based language**, nearly everything is an expression, and can be used in any context where an expression is expected

Try to be **as general as possible** in your code:

E.g., there is likely no need for a separate `compile_cond()` function, just call your `compile_expr()` function. If you need specific checks or operations (checking the type is `bool` or adding a `br` instruction), it should likely go into `compile_if()`. If you want to share condition-checking code between, say *if-then-else* and *while* loops, make a `compile_cond()` function, but have it call `compile_expr()` rather than reimplementing the same logic.

Avoid redundant code !

Keep It Simple

```
@.str = private unnamed_addr constant [14 x i8] c"Hello, world!\00", align 1
define i32 @main() #0 {
    %1 = alloca i32, align 4
    store i32 0, i32* %1, align 4
    %2 = call i32 @puts(i8* getelementptr inbounds ([14 x i8], [14 x i8]* @.str, i64 0, i64 0))
    ret i32 0
}
declare i32 @puts(i8* nocapture readonly) #1
attributes #0 = { noinline nounwind optnone uwtable
"correctly-rounded-divide-sqrt-fp-math"="false"
"disable-tail-calls"="false" "less-precise-fpmad"="false" "min-legal-vector-width"="0"
"no-frame-pointer-elim"="true" "no-frame-pointer-elim-non-leaf" "no-infs-fp-math"="false"
"no-jump-tables"="false" "no-nans-fp-math"="false" "no-signed-zeros-fp-math"="false"
"no-trapping-math"="false" "stack-protector-buffer-size"="8" "target-cpu"="x86-64"
"target-features"="+cx8,+fxsr,+mmx,+sse,+sse2,+x87" "unsafe-fp-math"="false"
"use-soft-float"="false" }
```

does (roughly) the same as the much simpler

```
@.str = constant [14 x i8] c"Hello, World!\00"
define i32 @main() {
    %1 = getelementptr [14 x i8], [14 x i8]* @.str, i64 0, i64 0
    %2 = call i32 @puts(i8* %1)
    ret i32 0
}
declare i32 @puts(i8*)
```


Generated Code vs Runtime

You can generate code inline for most operations, even in a dynamic language

This is **efficient** but:

- may lead to **code explosion**
- may be hard to do

You can **factorize** out common functionalities into a **library** provided with every VSOP program (called the **runtime**):

- It can be written in a different, usually lower-level language
- Some languages go as far as using a *virtual machine* to run programs, which are more data than code (e.g. Java)
- For VSOP, we will provide a runtime for I/O (the `Object` class)

How to Check/Implement the Object Class ?

Semantic analysis:

- The fact that Object exists and its method prototypes should be known during typechecking
- Add their *prototypes* manually to your symbol tables, or parse their definitions directly in VSOP (with dummy method bodies)

Code generation:

- **Don't generate** code for Object methods within your compiler !
- **Provide** the Object method definitions **separately**:
 - by **appending** their LLVM IR directly to generated IR
 - or by **linking** with an object file with their definitions (e.g. generated from C)

Alternatively, implement the *Foreign Function Interface* **FFI** extension and write the Object class directly in VSOP

Accessing the Runtime on the Platform

Your generated compiler is **not run** in your vsopcompiler folder, so you will not be able to access your runtime folder using **relative path** (`./runtime/`)

Solution: Use the Makefile to **copy** the runtime file(s) you need in a known location, e.g. `/usr/local/lib/vsop/`

⇒ Your compiler can access it using an **absolute path**

By-value vs By-reference

By-Value: Function arguments are **copied** before call, modifying an argument inside the callee **does not modify** corresponding variable at caller site

By-Reference: function arguments are **passed by pointer**, modifying an argument inside the callee **modifies caller variable**

By-Value is generally safer, modification allowed through explicit pointers

Like Java, VSOP uses by-value semantics for primitive types (like `int32`), but by-reference semantics for objects

Outline

The Assignment

The LLVM Language

Code Generation

Implementing the Dispatch

Q & A

Static vs Dynamic

- *Static properties* are known at **compile time**
- *Static behaviour* is what your compiler does **when compiling**
- **Safer** and more **efficient**
- *Dynamic properties* are only known at **runtime**
- *Dynamic behaviour* is what your compiled program does **when running**
- More **expressive**

Before implementing a feature, ask yourself whether it is **static**, **dynamic**, or if it needs **both** static and dynamic support

How to Implement Dispatch ?

How would you compile the following ?

```
class MyClass {  
  j : int32 <- 42;  
  someMethod(i : int32) : int32 {  
    i + j  
  }  
}  
  
// ...  
let myObject : MyClass <- new MyClass in  
myObject.someMethod(1942)
```

Simple Static Dispatch

Key idea: pass the **object** instance as **first argument** to methods.

```
// struct to keep fields data
typedef struct { int32_t j; } MyClass;
// Allocation and initialization
MyClass *MyClass_new(void) {
    MyClass *self = malloc(sizeof (MyClass));
    self->j = 42;
    return self;
}
// Methods
int32_t MyClass_someMethod(MyClass *self, int32_t i) {
    return i + self->j;
}

// ...
MyClass *myObject = MyClass_new();
MyClass_someMethod(myObject, 1942);
```


Field Inheritance

How to **reuse** parent field in child method ?

How to **add new** fields to children classes ?

```
class Parent {  
  i : int32 <- 42;  
  j : int32;  
}
```

```
class Child extends Parent {  
  k : int32;  
  sum() : int32 { i + j + k }  
}
```

```
// ...  
let child : Child <- new Child  
in child.sum()
```

Field Inheritance: Parent Class

Note that we decouple **allocation** and **initialization**

```
// Parent fields
typedef struct {
    int32_t i;
    int32_t j;
} Parent;

// Initialization of already allocated structure
void Parent_init(Parent *self) {
    self->i = 42;
    self->j = 0;
}

// Allocation and initialization
Parent *Parent_new(void) {
    Parent *self = malloc(sizeof (Parent));
    Parent_init(self); // Initialization
    return self;
}
```


Field Inheritance: Child Class

```
typedef struct {
    int32_t i; // FIRST: fields of parent, IN SAME ORDER
    int32_t j;
    int32_t k; // THEN: additional field(s) of Child
} Child;
// Initialization of child object
void Child_init(Child *self) {
    Parent_init((Parent *) self); // super()
    self->k = 0; }
// Allocation and initialization of a new child
Child *Child_new(void) {
    Child *self = malloc(sizeof (Child));
    Child_init(self); // Initialization
    return self; }
// Child method
int32_t Child_sum(Child *self) {
    return self->i + self->j + self->k; }

// ...
Child *child = Child_new();
Child_sum(child);
```

Field Inheritance: Memory Layout

```
;           i,    j  
%Parent = type { i32, i32 }  
;           i,    j,    k  
%Child  = type { i32, i32, i32 }
```



Method Inheritance

```
class Parent {  
  i : int32 <- 42;  
  inParent() : int32 { i }  
}  
  
class Child extends Parent {  
  inChild() : int32 { 1942 + i }  
}  
  
// ...  
let c : Child <- new Child in {  
  c.inParent();  
  c.inChild()  
}
```

Method Inheritance: Static Dispatch

Use a Child as a Parent using a cast

Same data layout for common fields of Child and Parent

```
typedef struct { int32_t i; } Parent;
void Parent_init(Parent *self) { self->i = 42; }
Parent *Parent_new(void) { /* ... */ }
int32_t Parent_inParent(Parent *self) { return self->i; }

typedef struct { int32_t i; /* Same as parent */ } Child;
void Child_init(Child *self) { /* ... */ }
Child *Child_new(void) { /* ... */ }
int32_t Child_inChild(Child *self) { return 1942 + self->i; }

// ...
Child *c = Child_new();
// c.inParent()
Parent_inParent((Parent *) c); // Cast needed (and OK) here
// c.inChild()
Child_inChild(c);
```

Static Dispatch is not enough

```
class Person {  
  i : int32;  
  name() : string { "Someone" }  
}
```

```
class John extends Person {  
  name() : string { "John" }  
}
```

```
class MyClass {  
  someMethod(p : Person): unit {  
    print(p.name()); // What to call here?  
    // Person_name(), John_name(), Mary_name(), ...?  
    ()  
  }  
}
```

Method Overriding: Function Pointers as Fields

```
typedef struct Person {
    int32_t i;
    char *(*_name)(struct Person *);
} Person;
char *Person_name(Person *self) { return "Someone"; }
void Person_init(Person *self) {
    self->i = 0;
    self->_name = &Person_name;
}

typedef struct John {
    int32_t i; // Same as Person
    char *(*_name)(struct John *); // John instead of Person, is it safe?
} John;
char *John_name(John *self) { return "John"; }
void John_init(John *self) {
    Person_init((Person *) self);
    self->_name = &John_name; // Overrides name()
}
// ...
p->_name(p) // p.name()
```


Method Overriding: Function Pointers as Fields

Using fields for methods works, but is very **wasteful**:

- **All objects** of the same class share the **same set** of methods
- But every object carries **one pointer per method**
- So ONE pointer PER method PER object $\Rightarrow$ **Waste of memory !**

Idea: **Share** the method pointers between objects of the same class

We need to **associate** each **object** with a **table of method pointers** (called the *vtable*):

- Either use **fat pointers**, i.e. a pointer to the object + a pointer to the *vtable*
- Or keep a pointer to the *vtable* in the object itself as a **field** (the *vtable* pointer)

Method Dispatch: Types

```
// Type for object instances
typedef struct {
    // Virtual function table (see below)
    MyClassVTable *vtable;
    // Fields (each object instance needs its own)
    int32_t j;
} MyClass;

// Type of the vtable
typedef struct {
    // Types of all the methods of MyClass
    void (*someMethod)(MyClass *, bool);
    int32_t (*someOtherMethod)(MyClass *, int32_t);
} MyClassVTable;
```

Method Dispatch: Constructor and Methods

```
void MyClass_init(MyClass *self) {  
    // Initialize fields, including virtual function table  
    self->vtable = &MyClass_vtable;  
    self->j = 42; // j : int32 <- 42  
}  
  
// someMethod(b : bool) : unit { ... }  
void MyClass_someMethod(MyClass *self, bool b) { /* ... */ }  
  
// someOtherMethod(i : int32) : int32 { i + j }  
int32_t MyClass_someOtherMethod(MyClass *self, int32_t i) {  
    return i + self->j;  
}  
  
// Actual function table object (only one needed)  
struct MyClassVTable MyClass_vtable = {  
    .someMethod = &MyClass_someMethod,  
    .someOtherMethod = &MyClass_someOtherMethod  
};
```

Method Dispatch: Calling a Method

```
// let myObject : MyClass <- new MyClass in  
MyClass *myObject = MyClass_new();  
// myObject.someOtherMethod(42)  
myObject->vtable->someOtherMethod(myObject, 42);
```

Note the **cost** of vtable-based dispatch: each method call now requires **two memory accesses** (*double dispatch*):

- One to get the vtable pointer
- One to get the method pointer

One can **optimize** out virtual calls when the actual type of the object is known (e.g. just after **new**)

In loops, one can also **cache** the method pointer

Method Dispatch: Adding Inheritance to the Mix

```
class Parent { (* ... *) }
class Child extends Parent { (* ... *) }

class X {
  someMethod(Parent p) : unit {
    // Static type of p is Parent, but dynamic type?
    p.overriddenMethod()
  }
}

void X_someMethod(X *self, Parent *p) {
  /* Will find the right method, whether p is a Parent,
   * a Child or any other descendant of Parent */
  p->vtable->overriddenMethod(p);
}
```

Method Dispatch: Adding Inheritance to the Mix (Cnt'd)

```
class Parent {  
  inheritedField : int32 <- 42;  
  inheritedMethod() : unit { (* ... *) }  
  overriddenMethod() : unit { (* ... *) }  
}
```

```
class Child extends Parent {  
  newField : bool <- true;  
  newMethod() : unit { (* ... *) }  
  overriddenMethod() : unit { (* ... *) }  
}
```

Method Dispatch: Adding Inheritance to the Mix (Cnt'd)

```
// A Child must be able to masquerade as a Parent
typedef struct {
    struct ChildVTable *vtable; // Virtual function table first
    int32_t inheritedField; // Parent fields, in the same order as parent!
    bool newField; // And, finally, Child's new fields
} Child;

// A ChildVTable must be able to masquerade as a ParentVTable
struct ChildVTable {
    // First, parent methods in the same order
    void (*inheritedMethod)(Child *), // Why not Parent * here?
    void (*overriddenMethod)(Child *),
    // Then child's new methods
    void (*newMethod)(Child *)
}

// Child VTable can mix inherited, overridden and new methods
struct ChildVTable Child_vtable = {
    // Necessary (but legit) cast for inherited method
    .inheritedMethod = (void (*)(Child *)) Parent_inheritedMethod,
    .overriddenMethod = Child_overriddenMethod,
    .newMethod = Child_newMethod
}
```


Method Dispatch: Adding Inheritance to the Mix (Cnt'd)

Chain constructors properly !

```
void Child_init(Child *self) {  
    Parent_init((Parent *) self); // Parent initializers  
    self->vtable = &Child_vtable; // Override vtable  
    self->newField = true;  
}
```

```
Child *Child_new() {  
    Child *self = malloc(sizeof (Child));  
    Child_init(self);  
    return self;  
}
```


In Summary



Beware of Symbol Conflicts

What if my class has a field named `vtable` ?

What if my class has a method called `init` ?

What if I have both:

- a class named `Base` with a method `jump_ship()` and
- a class named `Base_jump` with a method `ship()` ?

⇒ Symbol conflicts !

Use **name mangling** to avoid all possible conflicts in generated symbol names, e.g.:

- `_vtable` field (fields cannot start with `_`)
- `ClassName__methodName` (methods cannot begin with `_`)
- `ClassName__init` and `ClassName__vtable`

Outline

The Assignment

The LLVM Language

Code Generation

Implementing the Dispatch

Q & A

Questions and (possibly) answers

